

Mercedes-Benz Greener Manufacturing

March 15, 2021

0.0.1 Project 1 - Mercedes-Benz Greener Manufacturing

DESCRIPTION Reduce the time a Mercedes-Benz spends on the test bench.

0.0.2 Problem Statement Scenario:

Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include the passenger safety cell with the crumple zone, the airbag, and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium carmakers. Mercedes-Benz cars are leaders in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams.

To ensure the safety and reliability of every unique car configuration before they hit the road, Daimler's engineers have developed a robust testing system. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Daimler's production lines. However, optimizing the speed of their testing system for many possible feature combinations is complex and time-consuming without a powerful algorithmic approach.

You are required to reduce the time that cars spend on the test bench. Others will work with a dataset representing different permutations of features in a Mercedes-Benz car to predict the time it takes to pass testing. Optimal algorithms will contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Daimler's standards.

0.0.3 Requirements Analysis

0.0.4 Problem Summary

1. **GIVEN** we need to reduce the time that cars spend on the test bench
2. **WHEN** we apply Singular Value Decomposition to the dataset representing different permutations of features in a Mercedes-Benz car
3. **THEN** we can contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Daimler's standards.

Business Objectives and Constraints

- Predicting accurate time a car spends on the test bench
- No strict latency constraints, few seconds to a few minutes prediction time is okay, but not hours.

0.1 Step 1 : Load the Datasets

0.1.1 DataSet Overview

This dataset contains an anonymized set of variables, each representing a custom feature in a Mercedes car. For example, a variable could be 4WD, added air suspension, or a head-up display. The ground truth is labeled 'y' and represents the time (in seconds) that the car took to pass testing for each variable.

We have two comma-separated files: - **train.csv** — Contains the training set with 4209 rows (data points) and 378 columns (features) with labels - Here we can see there are 4209 datapoints indexing from 0 to 4208 and 378 columns/features. - We have three types of data in the train.csv dataset: > - float64(1): Dependent feature, testing time in seconds > - int64(369): Independent Binary features > - object(8): Independent Categorical features - **test.csv** — Contains the test set with 4209 rows (data points) and 377 columns (features) with no labels - Here we can see there are 4209 datapoints indexing from 0 to 4208 and 378 columns/features. - We have two types of data in the test.csv dataset: > - int64(369): Independent Binary features > - object(8): Independent Categorical features

We can see here we have the same number of data points in the train and test dataset.

```
[1]: import numpy as np
import pandas as pd
df_train = pd.read_csv('train.csv')
df_test = pd.read_csv('test.csv')
```

```
[2]: df_train.head()
```

```
[2]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8 ... X375 X376 X377 X378 X379 \
0    0  130.81  k  v  at  a  d  u  j  o ...    0    0    1    0    0
1    6   88.53  k  t  av  e  d  y  l  o ...    1    0    0    0    0
2    7   76.26 az  w   n  c  d  x  j  x ...    0    0    0    0    0
3    9   80.62 az  t   n  f  d  x  l  e ...    0    0    0    0    0
4   13   78.02 az  v   n  f  d  h  d  n ...    0    0    0    0    0
```

```
      X380 X382 X383 X384 X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
```

[5 rows x 378 columns]

```
[3]: df_test.head()
```

```
[3]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10 ... X375 X376 X377 X378 X379 X380 \
0    1 az  v   n  f  d  t  a  w    0 ...    0    0    0    1    0    0
```

1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

```
[4]: df_train.shape, df_test.shape
```

```
[4]: ((4209, 378), (4209, 377))
```

0.2 Step 2: Examine the data (for possible outliers)

- Before we do anything else, determine if there are any NaN values Or other irregularities

```
[5]: hasNans = df_train.isnull().sum().any()
hasDups = df_train.duplicated().sum().any()
hasDupIDs = df_train['ID'].duplicated().sum()>0
hasNans, hasDups, hasDupIDs
```


```
[5]: (False, False, False)
```

```
[6]: hasNans = df_test.isnull().sum().any()
hasDups = df_test.duplicated().sum().any()
hasDupIDs = df_test['ID'].duplicated().sum()>0
hasNans, hasDups, hasDupIDs
```


```
[6]: (False, False, False)
```

```
[7]: import matplotlib.pyplot as plt
from matplotlib import style
%matplotlib inline
plt.scatter(range(len(df_train)), np.sort(df_train.y.values), alpha=0.3)
```

```
[7]: <matplotlib.collections.PathCollection at 0x7fad884cbfa0>
```



```
[8]: #plot histogram
style.use('ggplot')
plt.figure(figsize=(7,7))
plt.hist(df_train.y,bins=100)
plt.xlabel('Independent Binary features ')
plt.ylabel('testing time in seconds')
plt.show()
```




Now remove the outliers values for the purposes of our model.

```
[9]: ulim = 190  
df_train.loc[df_train.y>ulim, 'y'] = ulim
```

```
[10]: import seaborn as sns  
sns.displot(df_train.y, bins=100, kde=False)
```

```
[10]: <seaborn.axisgrid.FacetGrid at 0x7fad78faeb20>
```



0.3 Step 3: Setup the `X_train`, `X_test`, `y_train`, `y_test` parameters for the models.

Identify any features from the train dataframe that are not part of the test dataframe

```
[11]: list(set(df_train.columns) - set(df_test.columns))
```

```
[11]: ['y']
```

- We conclude that the 'y' column of the train dataframe must be the item to be predicted.
- In effect, the 'y' column is the time, (in seconds) that is required to test a feature of a car on the assembly line.

```
[12]: y_train = df_train.y  
y_train.head()
```

```
[12]: 0    130.81
      1     88.53
      2     76.26
      3     80.62
      4     78.02
      Name: y, dtype: float64
```

Now remove the 'y' column from the train dataframe so that the train dataframe can be used to train the models.

```
[13]: X_train = df_train.drop(['y', 'ID'], axis = 1)
      X_train.head()
```

```
[13]:   X0 X1 X2 X3 X4 X5 X6 X8 X10 X11 ... X375 X376 X377 X378 X379 \
0    k v at a d u j o    0    0 ...    0    0    1    0    0
1    k t av e d y l o    0    0 ...    1    0    0    0    0
2   az w  n c d x j x    0    0 ...    0    0    0    0    0
3   az t  n f d x l e    0    0 ...    0    0    0    0    0
4   az v  n f d h d n    0    0 ...    0    0    0    0    0

      X380 X382 X383 X384 X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
```

[5 rows x 376 columns]

- We should use X_train and y_train from here
- Now what types of data are in the dataset?

```
[14]: print(X_train.dtypes.value_counts())
      print(df_test.dtypes.value_counts())
```

```
int64      368
object       8
dtype: int64
int64      369
object       8
dtype: int64
```

- Now remove any variables that just have a single unique value
- (as these do not add value to the model)

```
[15]: # check how many columns have a single unique value
      X_train_column_unique_values = X_train.apply(lambda x: pd.Series.nunique(x))

      # get the names of the columns to be removed
```

```

X_train_columns_remove = _
↪ X_train_column_unique_values[X_train_column_unique_values == 1].index.
↪ tolist()

# drop these columns from the dataframe
X_train = X_train.drop(X_train_columns_remove, axis = 1)
X_test = df_test.drop(X_train_columns_remove + ['ID'], axis = 1)

print("New shape of the train dataframe:", X_train.shape)
print("New shape of the train dataframe:", X_test.shape)

```

New shape of the train dataframe: (4209, 364)
New shape of the train dataframe: (4209, 364)

- Now remove any correlated variables

```
[16]: correlations = X_train.corr()
```

```

[17]: def remove_correlated(correlation_df):
    # list of all columns
    all_columns = correlation_df.columns
    chosen_columns = []
    removed_columns = []

    while len(all_columns) > 0:

        # choose the first column in the list
        col = all_columns[0]

        # add it to the chosen columns list
        chosen_columns.append(col)

        # set criteria to filter variables
        criteria = abs(correlation_df[col]) >= 0.6

        # get correlated variables except for the variable itself
        correlated_columns = list(set(correlation_df.loc[criteria, col].index) _
↪ set([col]))

        # reduce the overall columns to check
        all_columns = list(set(all_columns) - set(correlated_columns + [col]))

        # add columns to be removed in removed columns
        removed_columns.append(correlated_columns)

        # filter out removed variable from the correlation_df
        correlations_df = correlation_df[all_columns]

```



```
return chosen_columns
```

```
[18]: chosen_columns = remove_correlated(correlations)
```

```
[19]: X_train_final = X_train[chosen_columns]
X_train_final, X_test_final = X_train_final.align(X_test, join = 'inner', axis=
    ↳ 1)

print("Shape of train data:", X_train_final.shape)
print("Shape of test data:", X_test_final.shape)
```

Shape of train data: (4209, 185)

Shape of test data: (4209, 185)

```
[20]: # store feature names
feature_names = X_train_final.columns.tolist()

# convert to arrays
X_train_array = np.array(X_train_final)
y_train_array = np.array(y_train)
test_array = np.array(X_test_final)

from sklearn.model_selection import train_test_split
X_train_train, X_train_val, y_train_train, y_train_val =
    ↳ train_test_split(X_train_array, y_train_array, test_size = 0.2,
    ↳ random_state = 42)

print(X_train_train.shape)
print(X_train_val.shape)
print(y_train_train.shape)
print(y_train_val.shape)
```

(3367, 185)

(842, 185)

(3367,)

(842,)

0.4 Step 4: Model Building

0.4.1 sklearn.linear_model: Lasso

- Linear Model trained with L1 prior as regularizer (aka the Lasso)
- The optimization objective for Lasso is:
- $(1 / (2 * n_samples)) * ||y - Xw||^2_2 + \alpha * ||w||_1$
- Technically the Lasso model is optimizing the same objective function as the Elastic Net with `l1_ratio=1.0` (no L2 penalty).

```
[21]: from sklearn.linear_model import LassoCV
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score

# define the lasso cv model
cv_model = LassoCV(alphas = None, cv = 5, max_iter = 10000, random_state = 23)
cv_model.fit(X_train_train, y_train_train)
best_alpha = cv_model.alpha_

lasso_model = Lasso(alpha = best_alpha)
lasso_model.fit(X_train_train, y_train_train)

prediction = lasso_model.predict(X_train_val)
Lasso_score = r2_score(y_train_val, prediction)
Lasso_score
```

[21]: 0.5680120978533502

```
[22]: lasso_model = Lasso(alpha = best_alpha)
lasso_model.fit(X_train_array, y_train_array)
prediction = lasso_model.predict(test_array)
output = pd.DataFrame({'ID': df_test['ID'], 'y': prediction})
output.to_csv('sub_lasso_final.csv', index = False)
```

0.5 Numerical Data Selection

- **GIVEN** That only numerical data can be used by the models
- **WHEN** we select only the numerical columns
- **THEN** we can use the linear models with the datasets

```
[23]: X_train.head()
```

```
[23]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	\
0	k	v	at	a	d	u	j	o	0	0	...	0	0	1	0	0	
1	k	t	av	e	d	y	l	o	0	0	...	1	0	0	0	0	
2	az	w	n	c	d	x	j	x	0	0	...	0	0	0	0	0	
3	az	t	n	f	d	x	l	e	0	0	...	0	0	0	0	0	
4	az	v	n	f	d	h	d	n	0	0	...	0	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 364 columns]

```
[24]: y_train
```

```
[24]: 0      130.81
      1      88.53
      2      76.26
      3      80.62
      4      78.02
      ...
     4204    107.39
     4205    108.77
     4206    109.22
     4207     87.48
     4208    110.85
      Name: y, Length: 4209, dtype: float64
```

```
[25]: usable_columns = list(set(df_train.columns) - set(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8' ]))
      _y_train = df_train['y'].values
      _id_test = df_test['ID'].values
      _x_train = df_train[usable_columns]
      _x_test = df_test[usable_columns]
```

```
[26]: usable_columns = list(set(df_train.columns) - set(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8' ]))
      _y_train = y_train.values
      _id_test = df_test['ID'].values
      _x_train = _x_train[usable_columns]
      _x_test = df_test[usable_columns]
```

```
[27]: pd.set_option('mode.chained_assignment', None)
      for column in usable_columns:
          cardinality = len(np.unique(_x_train[column]))
          if cardinality == 1:
              _x_train.drop(column, axis=1)
              _x_test.drop(column, axis=1)
          if cardinality > 2:
              mapper = lambda x: sum([ord(digit) for digit in x])
              _x_train[column] = _x_train[column].apply(mapper)
              _x_test[column] = _x_test[column].apply(mapper)
```

```
[28]: X_train = _x_train
      y_train = _y_train
```

0.5.1 sklearn.model_selection: train_test_split

- Split arrays or matrices into random train and test subsets
- Quick utility that wraps input validation and `next(ShuffleSplit()).split(X,`

y)) and application to input data into a single call for splitting (and optionally subsampling) data in a oneliner.

- Read more in the [User Guide](#).

```
[29]: # store feature names
feature_names = X_train.columns.tolist()

# convert to arrays
X_train_array = np.array(X_train)
y_train_array = np.array(y_train)
test_array = np.array(X_test)

from sklearn.model_selection import train_test_split
X_train_train, X_train_val, y_train_train, y_train_val = train_test_split(
    X_train_array, y_train_array, test_size = 0.2,
    random_state = 42)

print(X_train_train.shape)
print(X_train_val.shape)
print(y_train_train.shape)
print(y_train_val.shape)
```

```
(3367, 368)
```

```
(842, 368)
```

```
(3367,)
```

```
(842,)
```

```
[30]: test_array = _x_train.values
```

0.5.2 sklearn.linear_model: Ridge

- Linear least squares with l2 regularization.
- Minimizes the objective function:
- $||y - Xw||^2_2 + \alpha * ||w||^2_2$
- This model solves a regression model where the loss function is the linear least squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when y is a 2d-array of shape (n_samples, n_targets)).

```
[31]: from sklearn.linear_model import RidgeCV
from sklearn.linear_model import Ridge

# define the lasso cv model
cv_model = RidgeCV(alphas = 10**np.linspace(10,-6,100)*0.5, cv = 5)
cv_model.fit(X_train_train, y_train_train)
best_alpha = cv_model.alpha_
```

```

# fit on entire training data
ridge_model = Ridge(alpha = best_alpha)
ridge_model.fit(X_train_train, y_train_train)

# get estimate on X_val
prediction = ridge_model.predict(X_train_val)
Ridge_score = r2_score(y_train_val, prediction)
Ridge_score

```

[31]: 0.5731494510022381

```

[48]: ridge_model = Ridge(alpha = best_alpha)
ridge_model.fit(X_train_array, y_train_array)
prediction = ridge_model.predict(test_array)
output = pd.DataFrame({'ID': df_test['ID'], 'y': prediction})
output.to_csv('sub_ridge_final.csv', index = False)

```

0.5.3 sklearn.model_selection: KFold

- LK-Folds cross-validator
- Provides train/test indices to split data in train/test sets. Split dataset into k consecutive folds (without shuffling by default).
- Each fold is then used once as a validation while the k - 1 remaining folds form the training set.

```

[49]: from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import KFold
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
%matplotlib inline

# hidden units = (100,), batch_size = 25, learning_rate = 0.00001, max_iter = 1000, score = 0.504
↪1000, score = 0.504

X_train_train, X_train_val, y_train_train, y_train_val = train_test_split(
    X_train_array, y_train_array, test_size = 0.2,
    random_state = 42)

mlp_model = MLPRegressor(activation = 'identity', solver = 'sgd', learning_rate_
↪= 'constant',
                        random_state = 42, learning_rate_init = 0.00001,
                        hidden_layer_sizes = (100,), max_iter = 1000,
↪batch_size = 25)

```

```

[50]: mlp_model.fit(X_train_train, y_train_train)
predictions = mlp_model.predict(X_train_val)

```

```
KFold_score = r2_score(y_train_val, predictions)
KFold_score
```

[50]: 0.5734967738018436

```
[35]: mlp_model.fit(X_train_array, y_train_array)
predictions = mlp_model.predict(test_array)
output = pd.DataFrame({'ID': df_test['ID'], 'y': predictions})
output.to_csv('sub_mlp.csv', index = False)
```

0.5.4 sklearn.linear_model: ElasticNet

- Linear regression with combined L1 and L2 priors as regularizer.
- Minimizes the objective function:
- $1 / (2 * n_samples) * ||y - Xw||^2_2$
- $\alpha * l1_ratio * ||w||_1$
- $0.5 * \alpha * (1 - l1_ratio) * ||w||^2_2$
- If you are interested in controlling the L1 and L2 penalty separately, keep in mind that this is equivalent to:
- $a * L1 + b * L2$
- where:
- $\alpha = a + b$ and $l1_ratio = a / (a + b)$
- The parameter `l1_ratio` corresponds to `alpha` in the `glmnet` R package while `alpha` corresponds to the `lambda` parameter in `glmnet`. Specifically, `l1_ratio = 1` is the lasso penalty. Currently, `l1_ratio <= 0.01` is not reliable, unless you supply your own sequence of `alpha`.

```
[36]: from sklearn.linear_model import ElasticNet
elas_model = ElasticNet(alpha=0.001, normalize=True)
elas_model.fit(X_train_train, y_train_train)
```

[36]: ElasticNet(alpha=0.001, normalize=True)

```
[37]: from math import sqrt
predictions1 = elas_model.predict(X_train_val)
Elas_score = elas_model.score(X_train_val, y_train_val)
Elas_score
```

[37]: 0.5278351276203694

```
[38]: elas_model.fit(X_train_array, y_train_array)
predictions2 = elas_model.predict(test_array)
Elas_score2 = elas_model.score(X_train_array, y_train_array)
output = pd.DataFrame({'ID': df_test['ID'], 'y': predictions2})
output.to_csv('sub_elas.csv', index = False)
Elas_score2
```

```
[38]: 0.5239550849260941
```

```
[39]: print(sqrt(mean_squared_error(y_train_val,predictions1)))  
      print(sqrt(mean_squared_error(y_train_array,predictions2)))
```

```
8.572772823779191
```

```
8.622657266500235
```

0.6 Step 7: Training using xgboost

XGBoost is an optimized distributed gradient boosting library designed to be highly efficient, flexible and portable. It implements machine learning algorithms under the Gradient Boosting framework. XGBoost provides a parallel tree boosting (also known as GBDT, GBM) that solve many data science problems in a fast and accurate way. The same code runs on major distributed environment (Hadoop, SGE, MPI) and can solve problems beyond billions of examples.

Parameters to be used for the xgboost model

```
[40]: X_train.head()
```

```
[40]:   X60  X77  X339  X75  X36  X376  X18  X190  X103  X245  ...  X359  X150  \  
0    0    0    0    0    0    0    1    0    0    0  ...    0    1  
1    0    0    0    0    0    0    1    0    0    0  ...    0    1  
2    0    0    0    1    0    0    0    0    0    0  ...    0    1  
3    0    0    0    0    0    0    0    0    0    0  ...    0    1  
4    0    0    0    0    0    0    0    0    0    0  ...    0    1
```

```
   X232  X86  X123  X250  X45  X80  X53  X65  
0    0    0    0    0    0    0    0    0  
1    0    0    0    1    0    1    0    0  
2    1    0    0    1    0    1    0    0  
3    1    0    0    1    0    1    0    0  
4    1    0    0    1    0    1    0    0
```

```
[5 rows x 368 columns]
```

```
[41]: X_test = _x_test  
      y_mean = y_train.mean()  
      print(X_train.shape)  
      print(X_test.shape)  
      print(y_train.shape)
```

```
(4209, 368)
```

```
(4209, 368)
```

```
(4209,)
```

```

[42]: import xgboost as xgb
from xgboost import XGBRegressor
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold

xgb_params = {
    'eta': 0.005,
    'max_depth': 3,
    'subsample': 0.95,
    'colsample_bytree': 0.6,
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'base_score': np.log(y_mean)
}

# form DMatrices for Xgboost training
#X_train_train, X_train_val, y_train_train, y_train_val
#dtrain = xgb.DMatrix(X_train_train, np.array(np.log(y_train_train)))
dtrain = xgb.DMatrix(X_train, np.array(np.log(y_train)))
dtest = xgb.DMatrix(X_test)

xgb_score = -1

# evaluation metric
def the_metric(y_pred, y):
    y_true = y.get_label()
    global xgb_score
    xgb_score = r2_score(y_true, y_pred)
    return 'r2', xgb_score

# xgboost, cross-validation
cv_result = xgb.cv(xgb_params,
                   dtrain,
                   num_boost_round=2000,
                   nfold = 3,
                   early_stopping_rounds=50,
                   feval=the_metric,
                   verbose_eval=True,
                   show_stdv=False
                   )

num_boost_rounds = len(cv_result)
print('num_boost_rounds=' + str(num_boost_rounds))

# train model
model = xgb.train(dict(xgb_params), dtrain, num_boost_round=num_boost_rounds)

```


[0]	train-rmse:0.12099	train-r2:0.00161	test-rmse:0.12099
	test-r2:0.00037		
[1]	train-rmse:0.12062	train-r2:0.00770	test-rmse:0.12062
	test-r2:0.00648		
[2]	train-rmse:0.12025	train-r2:0.01393	test-rmse:0.12024
	test-r2:0.01269		
[3]	train-rmse:0.11988	train-r2:0.01987	test-rmse:0.11988
	test-r2:0.01859		
[4]	train-rmse:0.11951	train-r2:0.02593	test-rmse:0.11951
	test-r2:0.02462		
[5]	train-rmse:0.11914	train-r2:0.03198	test-rmse:0.11914
	test-r2:0.03067		
[6]	train-rmse:0.11877	train-r2:0.03795	test-rmse:0.11877
	test-r2:0.03664		
[7]	train-rmse:0.11841	train-r2:0.04378	test-rmse:0.11841
	test-r2:0.04246		
[8]	train-rmse:0.11806	train-r2:0.04948	test-rmse:0.11806
	test-r2:0.04810		
[9]	train-rmse:0.11772	train-r2:0.05486	test-rmse:0.11773
	test-r2:0.05349		
[10]	train-rmse:0.11736	train-r2:0.06063	test-rmse:0.11737
	test-r2:0.05922		
[11]	train-rmse:0.11701	train-r2:0.06627	test-rmse:0.11702
	test-r2:0.06485		
[12]	train-rmse:0.11667	train-r2:0.07173	test-rmse:0.11668
	test-r2:0.07037		
[13]	train-rmse:0.11632	train-r2:0.07724	test-rmse:0.11633
	test-r2:0.07585		
[14]	train-rmse:0.11599	train-r2:0.08245	test-rmse:0.11600
	test-r2:0.08105		
[15]	train-rmse:0.11566	train-r2:0.08777	test-rmse:0.11566
	test-r2:0.08642		
[16]	train-rmse:0.11532	train-r2:0.09305	test-rmse:0.11533
	test-r2:0.09164		
[17]	train-rmse:0.11499	train-r2:0.09829	test-rmse:0.11500
	test-r2:0.09681		
[18]	train-rmse:0.11465	train-r2:0.10363	test-rmse:0.11467
	test-r2:0.10212		
[19]	train-rmse:0.11432	train-r2:0.10876	test-rmse:0.11434
	test-r2:0.10723		
[20]	train-rmse:0.11398	train-r2:0.11395	test-rmse:0.11401
	test-r2:0.11239		
[21]	train-rmse:0.11366	train-r2:0.11898	test-rmse:0.11369
	test-r2:0.11735		
[22]	train-rmse:0.11333	train-r2:0.12410	test-rmse:0.11336
	test-r2:0.12244		
[23]	train-rmse:0.11301	train-r2:0.12905	test-rmse:0.11304
	test-r2:0.12733		

[24]	train-rmse:0.11270	train-r2:0.13379	test-rmse:0.11274
	test-r2:0.13203		
[25]	train-rmse:0.11239	train-r2:0.13861	test-rmse:0.11243
	test-r2:0.13678		
[26]	train-rmse:0.11206	train-r2:0.14353	test-rmse:0.11211
	test-r2:0.14170		
[27]	train-rmse:0.11174	train-r2:0.14843	test-rmse:0.11179
	test-r2:0.14657		
[28]	train-rmse:0.11143	train-r2:0.15319	test-rmse:0.11148
	test-r2:0.15132		
[29]	train-rmse:0.11112	train-r2:0.15797	test-rmse:0.11117
	test-r2:0.15607		
[30]	train-rmse:0.11081	train-r2:0.16267	test-rmse:0.11086
	test-r2:0.16076		
[31]	train-rmse:0.11049	train-r2:0.16738	test-rmse:0.11055
	test-r2:0.16546		
[32]	train-rmse:0.11018	train-r2:0.17202	test-rmse:0.11024
	test-r2:0.17009		
[33]	train-rmse:0.10989	train-r2:0.17644	test-rmse:0.10995
	test-r2:0.17442		
[34]	train-rmse:0.10959	train-r2:0.18101	test-rmse:0.10965
	test-r2:0.17898		
[35]	train-rmse:0.10929	train-r2:0.18545	test-rmse:0.10935
	test-r2:0.18341		
[36]	train-rmse:0.10902	train-r2:0.18945	test-rmse:0.10908
	test-r2:0.18743		
[37]	train-rmse:0.10872	train-r2:0.19388	test-rmse:0.10878
	test-r2:0.19184		
[38]	train-rmse:0.10843	train-r2:0.19814	test-rmse:0.10850
	test-r2:0.19605		
[39]	train-rmse:0.10814	train-r2:0.20250	test-rmse:0.10821
	test-r2:0.20039		
[40]	train-rmse:0.10785	train-r2:0.20679	test-rmse:0.10792
	test-r2:0.20466		
[41]	train-rmse:0.10756	train-r2:0.21106	test-rmse:0.10763
	test-r2:0.20893		
[42]	train-rmse:0.10729	train-r2:0.21502	test-rmse:0.10736
	test-r2:0.21289		
[43]	train-rmse:0.10701	train-r2:0.21908	test-rmse:0.10709
	test-r2:0.21689		
[44]	train-rmse:0.10673	train-r2:0.22320	test-rmse:0.10680
	test-r2:0.22100		
[45]	train-rmse:0.10648	train-r2:0.22676	test-rmse:0.10656
	test-r2:0.22455		
[46]	train-rmse:0.10620	train-r2:0.23082	test-rmse:0.10628
	test-r2:0.22860		
[47]	train-rmse:0.10594	train-r2:0.23465	test-rmse:0.10602
	test-r2:0.23238		

```
[48]    train-rmse:0.10566      train-r2:0.23865      test-rmse:0.10574
test-r2:0.23637
[49]    train-rmse:0.10539      train-r2:0.24258      test-rmse:0.10547
test-r2:0.24030
[50]    train-rmse:0.10511      train-r2:0.24651      test-rmse:0.10520
test-r2:0.24420
num_boost_rounds=1
```

```
[43]: # Predict on train and test
y_train_pred = np.exp(model.predict(dtrain))
y_train_pred
```

```
[43]: array([100.62033 , 100.62033 , 100.52241 , ..., 100.705864, 100.62033 ,
        100.62033 ], dtype=float32)
```

```
[44]: # Predict on train and test
y_pred = np.exp(model.predict(dtest))
y_pred
```

```
[44]: array([100.52241 , 100.62033 , 100.52241 , ..., 100.62033 , 100.705864,
        100.62033 ], dtype=float32)
```

```
[45]: output = pd.DataFrame({'id': df_test['ID'].astype(np.int32), 'y': y_pred})
output.to_csv('sub_15_encoded.csv', index=False)
```

1 Summary

```
[46]: # X_train_train, X_train_val, y_train_train, y_train_val
```

```
[47]: print('R2 = {} with Elas'.format(Elas_score))
print('R2 = {} with KFold'.format(KFold_score))
print('R2 = {} with Ridge'.format(Ridge_score))
print('R2 = {} with Lasso'.format(Lasso_score))
print('R2 = {} with XGBoost'.format(xgb_score))
```

```
R2 = 0.5278351276203694 with Elas
R2 = 0.5734967738018436 with KFold
R2 = 0.5731494510022381 with Ridge
R2 = 0.5680120978533502 with Lasso
R2 = 0.2452701073023824 with XGBoost
```

1.1 Results

Out of the five models evaluated the Lasso model score showed it to be the most accurate to be used.

```
[ ]:
```

[]: